

Selly · AI Sales Coach

An AI-powered conversational coach for B2B sales reps

Live: b2b-sales-trainer.netlify.app

Stack: Claude API · Vanilla JS · Netlify Functions

Scope: Solo project · Instructional design, knowledge base research, system prompt engineering, frontend development, deployment

The Problem

Most sales training has an execution gap. Reps learn frameworks, watch recordings, pass quizzes, and still freeze on live calls because knowledge and reflex are different things. The reflex forms through repetition in conditions that resemble the real situation. Traditional options for that kind of practice, shadowing senior reps or learning on actual deals, either do not scale or are too expensive to get wrong.

Roleplay with a manager works, but it requires scheduling, availability, and the willingness to look incompetent in front of someone who influences your performance review. Most reps do not get enough of it, and when they do, the social stakes change what they are willing to try.

I built the B2B Sales Trainer to make deliberate practice available without the social overhead. A conversational coach that reacts to what you actually say, not what you meant to say, and tells you specifically what worked.

Meet Selly

The coach has a name and a face. Selly is a sketchy portrait illustration in monochromatic violet ink, designed to match the brand aesthetic and feel like someone who has heard every objection twice and is quietly entertained by the bad

ones. The name is a light piece of wordplay on “sell”, chosen deliberately to signal personality without taking itself too seriously.

Research on social presence in digital learning shows that a named, visually represented agent increases engagement and the sense of accountability during practice. This was an instructional design decision: a face to answer to changes how people show up to the conversation, including how much they are willing to expose what they do not know.

The Design Decisions

Why conversation and not a structured module. The core learning problem in sales is transfer of skill under pressure. Responding defensively to an objection is a habit; reciting LAER is a different cognitive task entirely. Deliberate practice addresses the habit. A conversational agent that reacts in real time, that pushes back when an answer is vague and adjusts when a point lands well, simulates the friction of a real call more than any sequence of screens can.

Why the selection screen asks about struggle areas, not modules. A rep knows they freeze on price objections. They may not know that falls under objection handling or that LAER is the relevant framework. Meeting users at the level of their actual experience, rather than asking them to navigate a curriculum, lowers the barrier to entry and produces more accurate self-assessment. The selection also primes attention: naming your struggle area before the session makes you more alert when that moment arrives.

Why there are two modes. Coach mode is the default: Selly pushes back, the user works for the answer, minimal scaffolding. Guide mode offers more structure before asking for a response. The toggle lives on the selection screen rather than in the chat because choosing a learning mode is a strategy decision, not a mid-session adjustment. Placing it upfront encourages deliberate practice rather than reactive switching when things get hard.

Why there is an “I’m not sure” option. A user who selects specific struggle areas has already done some self-assessment. Not every user arrives with that. The “I’m not sure” button bypasses the selection and instructs Selly to calibrate from zero: no assumed vocabulary, open diagnostic questions, no implied starting point. This makes the trainer accessible without changing the experience for someone who already knows what they need.

Why feedback quotes the user's own words. Abstract feedback is easy to agree with and easy to ignore. Selly is instructed to cite a specific phrase the user said, evaluate it concretely, and name one thing to do differently. The difference between “your answer could be more direct” and “when you said ‘I understand your concern’, that phrase signals you were listening but proves nothing” is the difference between feedback that lands and feedback that passes through.

The Knowledge Base

Before writing a line of the system prompt, I built a 34-page knowledge base from primary research across objection handling, discovery, buyer psychology, qualification frameworks, and value selling. Claims were cross-referenced across multiple sources before inclusion, statistics verified against their original methodology. Findings that changed design decisions were noted explicitly.

Two examples: the Gartner 2025 data on buying committees growing to nine to eleven stakeholders made the champion-versus-economic-buyer distinction central to the Reading the Buyer module, not peripheral. The evidence that 60% of buyers say sellers give discounts without being asked shaped how the value selling coaching frames discount discipline from the first exchange. The knowledge base was design input, not background reading.

What This Demonstrates

The trainer runs across five topic areas: Objection Handling, Discovery, Reading the Buyer, Timing and Qualification, and Value Selling. Users select struggle areas or indicate they are unsure, choose a mode, and optionally provide context about what they sell and who their typical buyer is. Selly uses that context to make roleplay scenarios specific: a skeptical VP at a mid-size SaaS company rather than a generic prospect. Memory between sessions is stored in localStorage and referenced naturally at the start of the next session.

The I chose conversational format because the goal is transfer, and transfer requires practice that actually resembles the real situation. I asked users to pick struggle areas instead of modules because meeting people at the vocabulary of their experience, not the curriculum's, produces better self-assessment and lowers the barrier to entry. Even though I initially designed this for reps who already had some sales experience, I added two modes after thinking about who else might use it: a single approach would exclude a lot of people who need this kind of practice most. And I instructed Selly to quote the user's own words in feedback because I've

seen enough learning happen in that one moment of recognition, when someone hears exactly what they said and why it didn't land.

The technical stack is deliberately simple: vanilla JavaScript, a Netlify Function as a backend proxy. The complexity in this project was instructional.

raquelmiranda.github.io · raquelgmiranda@gmail.com · linkedin.com/in/raquelgmiranda